

# Samuel Mucyo

sammucyo@college.harvard.edu  
sam-mucyo.github.io

## OBJECTIVE

---

Senior in Computer Science with coursework in high-performance computing, parallel programming, and systems optimization. Gained teaching experience as a CS50 Teaching Fellow, assisting students with programming and algorithmic concepts. Completed projects involving OpenMP, MPI, and profiling tools, building a foundation in parallel algorithms and performance analysis. Eager to support students in CS2050 while deepening my understanding of HPC techniques and architectures.

## EDUCATION

---

**Harvard University**, Cambridge, MA

Graduation Date: May 2025

Bachelor of Arts in Computer Science

GPA: 3.68/4.0

*Relevant Coursework:* High-Performance Computing (CS2050), Data Systems (CS1650), Systems Programming & Machine Organization (CS61), Data Structures & Algorithms (CS1240), Systems Development (CS107), Machine Learning (CS1810), Data Science I (CS109A)

## ACADEMIC EXPERIENCE

---

### Teaching Fellow, Introduction to Computer Science (CS50)

*Harvard University*

Fall 2022 and Fall 2023

- Conducted detailed code reviews for over 40 student projects and weekly problem sets through grading, providing feedback on code quality, memory management, and algorithm efficiency
- Mentored students during 6 weekly office hours, focusing on debugging issues across multiple programming languages: C, Python, and JavaScript
- Led weekly 2-hour hands-on lab sessions, covering topics such as introduction to data structures & algorithms in C, and web development with Flask, HTML, CSS, and JavaScript

### Parallelizing Urban Transit Construction with Minimum Spanning Trees

*Group Project: High Performance Computing for Science and Engineering (CS2050)* Spring 2024

- Collaborated in a team of five to implement parallel version of Kruskal's algorithm using OpenMP and OpenMPI for distributed computing, for final project
- Conducted comprehensive performance analysis using PAPI for hardware counter measurements
- Performed memory profiling using Valgrind's massif tool to optimize resource utilization
- Automated strong and weak scaling analysis through parameterized bash scripts

### Optimized Column-Store Database System

*Individual Project: Data Systems (CS1650)*

Fall 2024

- Developed a high-performance columnar database engine for efficient data retrieval and analysis (select-project-join) with optimized storage techniques, including B-trees for indexing and memory-mapped files for persistence
- Optimized query execution and concurrency by implementing multi-threaded scan operators, achieving a top 3 placement in class benchmarks for index operation and skewed data handling
- Implemented parallel batch select query execution using POSIX threads in C, with cache-conscious design and minimal data movement

### Astrolibrary

*Group Project: Systems Development for Computational Science (CS107)*

Fall 2023

- Followed the Software Engineering Development Life Cycle to deliver a Python library for astronomical spectral analysis
- Collaboratively designed API contracts, ensuring alignment with project requirements
- Developed and tested core functionality, supplemented with a comprehensive test suite using *pytest* for validation and documentation using *Sphinx*, enabling ease of use and future maintainability

- Integrated an automated CI/CD pipeline for seamless builds, testing, and deployment using GitHub Actions
- Coordinated with team members through regular meetings, version control, and code reviews

## INDUSTRY EXPERIENCE

---

### Amazon

Seattle, WA

*Software Development Engineer Intern at Amazon.com*

May 2024 - August 2024

- Created a reusable A/B testing framework for dynamic profile badges using Java/Spring MVC, reducing similar A/B experiment deployment time from weeks to days
- Improved API usage and accuracy by identifying and resolving misuse of the internal A/B API, which caused excessive triggers and signal noise; Conducted a statistical analysis to support the redesign

*Software Development Engineer Intern at Amazon.com*

May 2023 - August 2023

- Architected APIs for a rule-based risk monitoring and alarming system to reduce fraudulent activity detection time by working collaboratively with software engineers, data engineers, and investigators
- Developed a new Python-based backend with AWS CDK, enabling rapid and consistent deployments; leveraged AWS Lambda and DynamoDB for scalability and resource utilization

*Software Development Engineer Intern at AWS – Redshift*

June 2022 - August 2022

- Designed a serverless data lake using Python, AWS S3, and Glue to enhance bottleneck detection in Redshift infrastructure testing
- Automated data pipeline integration with internal visualization tool to provide detailed insights, enabling fast, data-driven decisions

## SKILLS

---

- **HPC & Systems Programming:** C, C++, POSIX threads, OpenMP, MPI, GDB, Valgrind
- **Cloud & Backend Development:** AWS (Lambda, S3, DynamoDB, CDK, Glue, Redshift), Java/Spring MVC, Python
- **Data Engineering:** SQL, database design, data pipelines, ETL, AWS Glue
- **Web Development:** JavaScript, React, Flask, HTML, CSS
- **Development Tools:** Git, Linux/Unix environments, CI/CD (GitHub Actions), Bash
- **Data Science & ML:** R, TensorFlow, Pandas, statistical analysis